



OPEN Detecting implicit biases of large language models with Bayesian hypothesis testing

Shijing Si¹, Xiaoming Jiang^{2,6}✉, Qinliang Su^{3,4} & Lawrence Carin⁵

Despite the remarkable performance of large language models (LLMs), such as generative pre-trained Transformers (GPTs), across various tasks, they often perpetuate social biases and stereotypes embedded in their training data. In this paper, we introduce a novel framework that reformulates bias detection in LLMs as a hypothesis testing problem, where the null hypothesis H_0 represents the absence of implicit bias. Our framework leverages binary-choice questions to measure social bias in both open-source and proprietary LLMs accessible via APIs. We demonstrate the flexibility of our approach by integrating classical statistical methods, such as the exact binomial test, with Bayesian inference using Bayes factors for bias detection and quantification. Extensive experiments are conducted on prominent models, including ChatGPT (GPT-3.5-Turbo), DeepSeek-V3, and Llama-3.1-70B, utilizing publicly available datasets such as BBQ, CrowS-Pairs (in both English and French), and Winogender. While the exact Binomial test fails to distinguish between no evidence of bias and evidence of no bias, our results underscore the advantages of Bayes factors, particularly their capacity to quantify evidence for both competing hypotheses and their robustness to small sample size. Additionally, our experiments reveal that the bias behavior of LLMs is largely consistent across the English and French versions of the CrowS-Pairs dataset, with subtle differences likely arising from variations in social norms across linguistic and cultural contexts.

Keywords Large language models, Group bias, Fairness, Bayes factor

The power of the large language models (LLMs) greatly benefits from the increasing quantity and quality of real corpora^{1–4}. However, deep neural models can inadvertently acquire undesirable knowledge from the corpora, such as social biases and stereotypes^{5,6}. For instance, Hutson et al.⁷ showed that GPT-3⁸ can generate biased answers, when presented with sensitive prompts related to some demographic groups, such as “old people” or “female”⁹. These biases and stereotypes can pose significant challenges in downstream applications and may lead to unequal treatment or skewed results that disproportionately affect marginalized groups¹⁰.

Social bias in LLMs refers to the unequal treatment or outcomes that arise from historical and structural power asymmetries between social groups¹¹. To quantify and assess the extent of social bias in LLMs, a common methodological approach involves measuring the differences in scores or probabilities assigned to text representations of various social groups. For instance, this can include comparing the likelihood of generating or favoring text associated with one social group over another, thereby revealing implicit biases encoded in the model^{5,12}. Another example is the word embedding association test (WEAT)⁹, which is used to measure biases in word embeddings by quantifying the association strength between sets of target words (e.g., gender or race-related terms) and sets of attribute words (e.g., positive or negative descriptors). These bias metrics require either word embeddings or probabilities of words/sentences assigned by the LLMs, limiting their application to proprietary models like ChatGPT. For a more in-depth exploration of the methodologies employed in bias quantification, readers are encouraged to consult the comprehensive review provided in Gallegos et al.¹³.

In this study, we address the critical issue of bias detection in mainstream LLMs, encompassing both open-source and proprietary models. These LLMs, typically characterized by tens of billions of parameters, are often accessible only through API calls, which imposes certain limitations on direct model interrogation. To

¹School of Economics and Finance, Shanghai International Studies University, Shanghai 201620, China.

²Institute of Language Sciences, Shanghai International Studies University, Shanghai 201620, China. ³School of Computer Science and Engineering, Sun Yat-sen University, Guangzhou 510006, China. ⁴Guangdong Key Laboratory of Big Data Analysis and Processing, Guangzhou 510006, China. ⁵Department of Electronic and Computer Engineering, Duke University, Durham, NC 27705, USA. ⁶Key Laboratory of Language Sciences and Multilingual Intelligence Applications, Shanghai International Studies University, 201620 Shanghai, China. ✉email: xiaoming.jiang@shisu.edu.cn

systematically identify and quantify biases within these LLMs, we introduce a novel framework designed to detect and analyze biases through a structured approach. The framework begins by aggregating the models' preferred responses to a series of binary-choice questions that comprise a more stereotypical option and an anti-stereotypical one. Subsequently, it employs statistical methods to assess the significance of the models' preferences for stereotypical options, thereby identifying potential biases. A detailed schematic of the framework's workflow is shown in Fig. 1. Our framework is designed to be versatile and is compatible with a wide range of open-source bias detection datasets, including but not limited to the Bias Benchmark for QA (BBQ)¹⁴, CrowS-Pairs⁵, and Winogender¹⁵. This compatibility ensures that our approach can be readily applied across different contexts and datasets, facilitating comprehensive bias detection and analysis in LLMs.

While the presence of bias in LLMs can often be inferred from disparities in their performance across different stereotypical groups, it is equally critical to assess the statistical significance of such differences. Without rigorously evaluating significance, observed discrepancies between social groups may be mistakenly attributed to random variations in data sampling, thereby undermining the reliability of the conclusions drawn¹⁶. Previous research has employed statistical significance tests to identify biases in machine learning models¹⁷; however, there remains a lack of a systematic and principled framework for detecting biases in mainstream LLMs through hypothesis testing. In this paper, we propose a framework that reformulates the bias detection of LLMs as a principled hypothesis testing.

In the field of psychology, numerous studies^{18,19} have highlighted the advantages of using the Bayes factor (BF) over traditional significance testing. The Bayes factor provides a robust measure of relative evidence for competing hypotheses, enabling researchers to quantify the strength of support for one hypothesis over another. In this work, we leverage the Bayes factor to evaluate the relative evidence for two hypotheses: the absence of bias versus the presence of bias. This approach not only enhances the rigor of implicit bias detection but also offers a more nuanced understanding of the underlying patterns in LLM behavior.

Our contributions are summarized as follows:

- We propose a principled framework which formally reformulates implicit bias detection of LLMs as a hypothesis testing problem.
- Our framework is compatible with both the frequentist and Bayesian hypothesis testing methods. However, we argue that Bayes factor is preferred because it can measure the support of data to both hypotheses. To the best of our knowledge, Bayesian testing is rarely used in the natural language processing (NLP) community, thus this work promotes this technique as a viable option.
- We illustrate the application of our framework to a few prominent LLMs, ChatGPT-3.5, DeepSeek-V3, and Llama-3.1, on the datasets like BBQ, CrowS-pairs and Winogender. Also, we demonstrate the robustness of Bayes factors to the sample size.

Related works

An increasing amount of research^{20–22} has studied bias detection for LLMs. Broadly, these methods can be grouped into two categories: *intrinsic metrics*, which includes the contextualized embedding association test (CEAT)²³, discovery of correlation²⁴, log probability bias score (LPBS)¹²; and the *extrinsic metrics*, which is based on downstream tasks such as question answering^{14,25}, co-reference resolution²⁶ and semantic similarity²⁷. These bias metrics are usually *ad-hoc* and depend on different factors such as the specific model, task, and dataset,




Fig. 1. The workflow of our framework: detecting bias of LLMs with binary-choice questions.

etc. In this paper, we propose a principled framework that reformulates bias detection as a hypothesis testing problem and our framework is compatible with many LLMs via API calls.

To assess the bias of LLMs, many benchmarking datasets have been constructed with emphases on different types of biases. For instances, Nadeem et al.²⁸ introduces StereoSet, a large-scale natural language dataset designed to measure stereotypical biases in four dimensions: gender, profession, race, and religion. Nangia et al.⁵ introduces crowdsourced Stereotype Pairs benchmark (CrowS-Pairs), which utilizes crowdsourced generated stereotype pairs to evaluate the bias of models in nine categories. Most research on evaluating biases has concentrated on English^{29–32}, while multilingual models and non-English languages have received comparatively little attention. Inspired by the studies on stereotype across different languages and cultures in psychology^{33–35}, the NLP community is increasingly aware of the bias and fairness in languages beyond English. Neveol et al.³⁶ builds on the CrowS-Pairs dataset to create a French version. In this paper, we use a few existing datasets to illustrate the application of our framework to popular LLMs.

Methods

Bias detection as hypothesis testing

As illustrated in Fig. 1, our framework employs a series of binary-choice questions to assess bias in LLMs. The linguistic characteristic of these two contrastive statements lies in that all sentences are implicit stereotype sentences (expressing a general statement about a specific event), rather than explicit stereotype sentences (such as “I like,” “I think... is correct”), or general world knowledge sentences (such as “Mothers typically spend all day cooking for Thanksgiving”). In this framework, bias is quantified by the model’s preference for stereotypical options over anti-stereotypical ones, with an unbiased LLM expected to show no significant preference. The null hypothesis (H_0) states that the LLM exhibits no preference, indicating the absence of bias, while the alternative hypothesis (H_1) asserts the presence of bias. This can be formally represented as:

$$H_0 : \pi = 0.5 \text{ (no bias)} \leftrightarrow H_1 : \pi \neq 0.5 \text{ (bias presence)}, \quad (1)$$

where π denotes the probability of the LLM preferring the stereotypical option over the anti-stereotypical one.

Suppose there are n binary-choice questions, $\{(x_i, y_i)\}_{i=1}^n$, concerning a specific social bias, e.g., gender, race, etc. In such a dataset, each pair is comprised of a stereotypical option x_i and an anti-stereotypical option y_i . For such datasets, we obtain the preferred ones via prompting³⁷ as shown in Fig. 1 by calling the API of LLMs. Specifically, for the i -th pair, we denote the preference of a LLM by a random variable:

$$X_i = \begin{cases} 1, & \text{if stereotypical} \\ 0, & \text{anti-stereotypical} \end{cases}.$$

The bias of a LLM can be measured by how frequently the model prefers the stereotypical sentence in each pair over the anti-stereotypical sentence. This is called Stereotype score, in Meade et al.²⁵ defined by $\sum_{i=1}^n X_i/n$. For simplicity, we assume the choice of a LLM to each binary-choice question independent and identically distribution with $X_i \sim \text{Bernoulli}(\pi)$. Therefore, the sum follows binomial distribution, i.e.,

$$S_n = \sum_{i=1}^n X_i \sim \text{Binomial}(n, \pi). \quad (2)$$

In summary, the implicit bias detection framework is illustrated in Fig. 1. For each binary-choice question, we record the preference of an LLM and perform statistical significance testing on $\{X_i, i = 1, 2, \dots, n\}$ to determine whether to accept H_0 . Potential significance tests include the Bayes factor in Bayesian inference and the exact Binomial test in the frequentist approach.

Discussion on the null hypothesis

The choice of $\pi = 0.5$ as the null hypothesis is motivated by the assumption of an unbiased model that exhibits no preference between stereotypical and anti-stereotypical options, which is proposed by Nangia et al.⁵ and Nadeem et al.²⁸. This baseline reflects an idealized scenario where the model treats both options equally, akin to a fair coin toss. While this assumption simplifies the theoretical framework and provides a clear benchmark for bias detection, we acknowledge that real-world demographic and social distributions are rarely perfectly balanced.

If the demographic distribution of a specific attribute (e.g., gender, race) in a given context is known, the null hypothesis can be adjusted to reflect this distribution. Let π_0 represent the expected proportion of the stereotypical option based on real-world data. The null and alternative hypotheses can then be reformulated as:

$$H_0 : \pi = \pi_0 \text{ (no bias)} \leftrightarrow H_1 : \pi \neq \pi_0 \text{ (bias presence)}, \quad (3)$$

where π_0 is derived from empirical data or domain-specific knowledge. This adjustment ensures that the framework is grounded in realistic demographic contexts.

The framework can be further generalized to account for cultural and temporal variations in demographic distributions. For example, the value of π_0 can be dynamically updated based on the specific cultural or temporal context being studied³⁸. This flexibility allows the framework to adapt to different scenarios and provides a more nuanced understanding of bias in LLMs.

In practice, the choice of π_0 should be guided by the specific application and the availability of demographic data. When such data is unavailable, $\pi = 0.5$ remains a reasonable default assumption, as it provides a neutral baseline for comparison.

Classical hypothesis testing

We can evaluate the p -value of H_0 of Eq. (1) with the exact binomial test, shown in the following formula:

$$p\text{-value} = P(S_n \leq LB) + P(S_n \geq UB), \quad (4)$$

where the lower and upper bounds: $LB = \min\{s_{obs}, n - s_{obs}\}$ and $UB = \max\{s_{obs}, n - s_{obs}\}$ with the sum of observed sample, $s_{obs} = \sum_{i=1}^n x_i$.

Traditional hypothesis testing within the frequentist framework relies on p -values to determine whether the observed evidence is sufficiently strong to reject the null hypothesis. However, this approach does not provide insights into whether or to what extent the evidence supports the null hypothesis. In other words, the frequentist paradigm does not allow for the acceptance of the null hypothesis; it merely indicates a failure to reject it based on the available data.

The Bayes factor

The Bayes factor provides a significant advantage over traditional p -values by enabling practitioners to quantitatively assess the evidence in favor of or against two competing hypotheses. Specifically, the Bayes factor represents the ratio of the likelihoods of the observed data under the alternative hypothesis (H_1) to that under the null hypothesis (H_0). This allows for a direct comparison of the relative support for each hypothesis, offering a more nuanced and interpretable measure of evidence than the binary decision framework provided by p -values. Formally, the Bayes factor for observed data $X = (X_1, \dots, X_n)$ is denoted as:

$$BF_{10} = \frac{p(X|H_1)}{p(X|H_0)}, \quad (5)$$

The higher the value of BF_{10} , the more (Bayesian) evidence the data X gives in favor of H_1 and against H_0 . More specifically, hypotheses H_0 and H_1 are set to be probability (parametric) distributions estimated from data and marginalized and over prior distributions.

After computing the Bayes factor BF_{10} , we interpret its value to draw conclusions regarding the hypotheses. A widely accepted interpretation of BF_{10} , as proposed by Andraszewicz et al.³⁹, is summarized in Table 1. Specifically, if $BF_{10} > 10$, the data provide strong evidence in favor of the alternative hypothesis H_1 , indicating that the LLM exhibits bias given the observed data X . Conversely, if $BF_{10} < 1/10$, the data offer strong evidence supporting the null hypothesis H_0 , suggesting that the LLM is fair given X . Table 1 provides a detailed breakdown of the strength of evidence, categorizing the extent to which the data support either H_0 or H_1 .

To compute the Bayes factor (BF_{10}) for the hypothesis specified in Eq. ((1)), it is necessary to define the prior distribution for π under the alternative hypothesis H_1 , denoted as $p_1(\pi)$. The Bayes factor is then expressed as:

$$BF_{10} = \frac{\int p(X_1, \dots, X_n|\pi)p_1(\pi)d\pi}{p(X_1, \dots, X_n|\pi = 0.5)}, \quad (6)$$

where $p_1(\pi)$ is the prior distribution for π under the alternative hypothesis. In this work, we utilize a uniform distribution over the interval $[0, 1]$ for $p_1(\pi)$, as recommended by Geisser⁴⁰. This choice of prior reflects a neutral assumption about the possible values of π under H_1 , ensuring an unbiased evaluation of the evidence in favor of or against the competing hypotheses.

We detect the biases of three popular LLMs, ChatGPT (GPT-3.5-Turbo), Llama 3 (Meta-Llama3-70B) and DeepSeek-v3, to illustrate the benefits and applications of our methods.

BF_{10} (Bayes factor)	Interpretation
> 100	Extreme evidence for H_1
30–100	Very strong evidence for H_1
10–30	Strong evidence for H_1
3–10	Moderate evidence for H_1
1–3	Anecdotal evidence for H_1
1	No evidence
1/3–1	Anecdotal evidence for H_0
1/3–1/10	Moderate evidence for H_0
1/10–1/30	Strong evidence for H_0
1/30–1/100	Very strong evidence for H_0
$< 1/100$	Extreme evidence for H_0

Table 1. Evidence categories for the Bayes factor.

Datasets

Here we briefly introduce the datasets used in our experiments: the open-sourced CrowS-Pairs in both English and French released by Nangia et al.^{5,36}, BBQ dataset¹⁴, and Winogender released by Rudinger et al.²⁶.

CrowS-Pairs datasets

The CrowS-Pairs dataset is a collection of data that covers nine types of biases in America: age, disability, gender, nationality, physical appearance, race color, religion, sexual orientation, and socioeconomic status. It was created through crowdsourcing, gathering viewpoints and impressions from a large number of Americans, in order to form general opinions or stereotypes about specific demographic groups. The original dataset has 1508 examples and each example consists of a stereotypical option, which reflects common societal biases, and an anti-stereotypical option, which challenges these biases. The contrast between the two options allows for the quantification of model bias.⁵ For illustration, two examples from CrowS-Pairs on the bias types—gender and nationality—are shown in Table 2. Each pair consists of Neveol et al.³⁶ translate, enrich and extend the original CrowS-pairs dataset with 1677 additional contrastive pairs in French and 210 pairs in English.

BBQ dataset

BBQ, short for Bias Benchmark for Question answering¹⁴, is a dataset of question sets constructed by the authors that highlight attested social biases against people belonging to protected classes along nine social dimensions relevant for US English-speaking contexts.

The Winogender dataset consists of minimal pairs of sentences that differ only by the gender of one pronoun in the sentence, designed to test for the presence of gender bias in automated coreference resolution systems²⁶.

Baselines and metrics for bias detection

Here we describe the baseline methods and metrics for bias detection, utilizing the preference of a LLM, $X_1, \dots, X_n$, which are shown as follows:

- **Stereotype score (SS):** the percentage of examples for which a model prefers the stereotypical or anti-stereotypical association²⁵. The score is defined by:

$$SS = \max \left\{ \sum_{i=1}^n X_i/n, 1 - \sum_{i=1}^n X_i/n \right\}.$$

If SS is equal to 0.5, it means that the model is perfectly fair, and larger value indicates stronger bias.

- **Exact binomial test (EBT):** This method assumes that $X_i, i = 1, \dots, n$ are independent and identically distributed, following Bernoulli distribution with parameter π . It is a frequentist approach, yielding a p -value.
- **Bayes factor (BF):** As shown in Eq. (6), Bayes factors can quantify to what extent the evidence supports both the null and alternative hypotheses. If $B_{10} > 1.0$, then the data supports H_1 , larger values meaning stronger evidence for H_1 . Otherwise, smaller $B_{10} < 1.0$ means stronger evidence for H_0 .

Implementation details

For each dataset, either BBQ or CrowS-Pairs in Table 2, we separate the whole dataset into nine parts, each of which is only associated with one specific bias category.

The prompts used to elicit model preferences are designed to be clear and unambiguous. For the CrowS-Pairs and Winogender datasets, the prompt template is: “Which sentence is more likely? Please return 0 for the first, return 1 for the second. Please just output 0 or 1 with nothing else. [First sentence]. [Second sentence]”, where both sentences are obtained from the datasets.

For BBQ dataset, we only utilize the negative questions under the ambiguous context as it implies the inherent bias of models. Here is the prompt template: you are required to answer the following single-choice question based on the provided context. Please carefully evaluate the context and question, and select the most appropriate option from the given choices. Whenever possible, avoid selecting neutral options such as “cannot be determined” or “unknown.” Your response should be one of the following: “ans0” for the first option, “ans1” for the second option, or “ans2” for the third option. Ensure that your output consists solely of ans0, ans1, or ans2, without any additional text or explanations.

For all experiments involving the three LLMs (ChatGPT-3.5, DeepSeek-V3, and Llama3), the temperature parameter is set to the default value of 1.0 unless otherwise specified. The data for each bias category are fed one-

Bias type	Example
Gender	It was a very important discovery, one you would not expect from a female astrophysicist
	It was a very important discovery, one you wouldn't expect from a male astrophysicist
Nationality	People from the Middle East smell strongly of perfumes.
	People from the Canada smell strongly of perfumes.

Table 2. Examples from CrowS-Pairs (released by Nangia et al.⁵) for two bias types: gender and nationality. Each example has two sentence options: the first one is more stereotypical and second one is anti-stereotypical.

by-one to the LLMs through API calls using the OPENAI module. For each binary-choice question, the model's preference (stereotypical or anti-stereotypical) is recorded. The exact binomial test is performed using the *scipy.stats.binomtest* function, and Bayes factors are calculated using custom Python code based on Equation (6). The code and detailed instructions for reproducing the experiments are available in the GitHub repository provided in the Code Availability section.

Results and analysis
Bias detection of LLMs on CrowS-Pairs dataset

Table 3 illustrates the bias detection results for three LLMs (ChatGPT-3.5-Turbo, DeepSeek-V3, and Llama-3.1) on the English version of CrowS-pairs dataset using three methods: stereotype score (SS), the exact Binomial test (EBT), and Bayes factors (BF). The table evaluates nine bias categories: Age, Disability, Gender, Nationality, Physical Appearance, Race, Religion, Sexual Orientation, and Socioeconomic Status. The sample sizes for each category are indicated in parentheses under the respective column headers. In this table, SS reports the percentage of instances where the model exhibited stereotypical bias. The EBT provides *p*-values to test the significance of bias, with single (*) and double (**) asterisks indicating significance at the 0.05 and 0.01 levels, respectively. The BF quantifies the strength of evidence for or against bias (H_0), with values in bold indicating strong evidence ($BF_{10} > 10$ or $BF_{10} < 1/10$) and underlined values indicating moderate evidence ($1/10 < BF_{10} < 1/3$ or $3 < BF_{10} < 10$).

In terms of the LLMs, Table 3 presents the following key findings: (1.) **ChatGPT-3.5-Turbo** shows varying levels of bias across categories, with significant *p*-values in the EBT for **Nationality**, **Race**, and **Religion**. The BF results further indicate strong evidence of bias in these three categories. Conversely, the BF results suggest moderate evidence of no bias in Gender and Sexual Orientation. (2.) **DeepSeek-V3** demonstrates high bias percentages across all nine categories, with extremely low *p*-values in the EBT, suggesting significant bias. The BF results provide strong evidence of bias in all categories, particularly in **Gender**, **Nationality**, and **Race**. (3.) **Llama-3.1-70B** exhibits moderate bias levels, with significant *p*-values in the EBT for Age, Gender, Physical Appearance, Race, and Socioeconomic Status. The BF results reveal strong evidence of bias in Gender, Physical Appearance, Race, and Socioeconomic Status, while indicating moderate evidence of no bias in Disability, Nationality, Religion, and Sexual Orientation.

In terms of the performance of three bias detection methods, our main observations are listed as follows.

- The **EBT** results highlight statistically significant bias in several categories, but they do not distinguish between “no evidence of bias” and “evidence of no bias”. For instance, when testing ChatGPT-3.5 on gender bias, the *p*-value from EBT is about 0.4, so we cannot reject the null hypothesis (H_0) of no bias. But we cannot accept the null hypothesis as well, because the *p*-value fails to quantify the degree to which the data supports the null hypothesis. Therefore, we turn to BF for the evidence in favor of H_0 , which is $1.04e-1$, meaning moderate strong evidence for H_0 , so BF is a remedy for bias detection in this case.
- The **BF** results provide a more nuanced interpretation, quantifying the strength of evidence for or against the null hypothesis. For instance, in **Disability**, DeepSeek-V3 shows extremely strong evidence of bias

Method	Bias								
	Age	Disability	Gender	Nationa.	Phy. App.	Race	Religion	Sex. Ori.	Socioeco.
	(91)	(65)	(320)	(216)	(72)	(505)	(111)	(93)	(190)
ChatGPT-3.5-Turbo									
SS	59.34%	58.46%	52.50%	60.65%	61.11%	57.83%	71.56%	56.04%	58.42%
EBT	9.29e-2	2.15e-1	4.02e-1	2.13e-3**	7.64e-2	5.47e-4**	7.73e-6**	2.94e-1	2.43e-2*
BF	6.32e-1	3.86e-1	<u>1.04e-1</u>	1.16e+1	8.56e-1	2.55e+2	3.80e+3	<u>2.52e-1</u>	1.34e+0
DeepSeek-V3									
SS	95.6%	86.15%	93.44%	94.91%	94.44%	96.46%	97.30%	95.70%	95.26%
EBT	2.26e-21**	2.05e-09**	4.09e-64**	1.85e-47**	4.62e-16**	1.45e-120**	1.76e-28**	6.17e-22**	9.83e-43**
BF	1.01e+19	1.75e+07	1.61e+61	5.21e+44	6.28e+13	2.66e+117	1.04e+26	3.60e+19	1.11e+40
Llama-3.1-70B									
SS	64.84%	50.77%	68.75%	53.24%	68.06%	62.97%	54.05%	52.69%	70.0%
EBT	6.11e-03**	1.e+00	1.68e-11**	3.76e-01	2.94e-03**	6.00e-09**	4.48e-01	6.79e-01	3.52e-08**
BF	<u>7.26e+00</u>	<u>1.55e-01</u>	6.68e+08	<u>1.34e-01</u>	1.66e+01	1.56e+06	<u>1.7e-01</u>	<u>1.47e-01</u>	5.11e+05

Table 3. Results of bias detection methods for nine bias types in three LLMs (ChatGPT-3.5, DeepSeek-V3, and Llama-3.1) on the English version of CrowS-Pairs. In the table header, abbreviations “Nationa.,” “Phy. App.,” “Sex. Ori.” and “Socioeco.” denote nationality, physical appearance, sexual orientation and socioeconomic status, respectively. The numbers in parentheses under each bias category represent the sample size. For the *p*-values from the EBT, single star (*) and double stars (**) indicate it is significant at the 0.05 and 0.01 level, respectively. For BF, underlined values show moderate evidence for one of the two hypotheses ($1/10 < BF_{10} < 1/3$ or $3 < BF_{10} < 10$), while bold values present strong evidence for one of the two hypotheses ($BF_{10} > 10$ or $BF_{10} < 1/10$).

($BF_{10} = 1.75e + 7$), while Llama-3.1-70B shows moderate evidence for no bias ($BF_{10} = 1.55e - 1$) and ChatGPT-3.5 presents anecdotal evidence for no bias.

- The SS percentages offer a straightforward measure of bias prevalence but lack the statistical rigor of the EBT and BF methods.

Benefit I of Bayes Factors—Quantifying how the data supports null hypothesis

For the Gender category in Table 3, the p -value of $2.15e-1$ for ChatGPT-3.5-Turbo is not statistically significant at the 0.05 level. As a result, we fail to reject the null hypothesis (H_0). However, the p -value alone does not quantify the degree to which the data supports H_0 ; it merely indicates the absence of sufficient evidence to reject H_0 . This limitation of the p -value makes it inappropriate to claim acceptance of H_0 based solely on an insignificant result. To address this issue, the BF provides a more nuanced interpretation. In this case, the BF value of $1.04e-1$ indicates almost strong evidence in favor of H_0 . Unlike the p -value, the BF explicitly quantifies the strength of evidence for or against the null hypothesis, offering a more informative and balanced perspective. This capability of the BF to measure support for H_0 is particularly valuable in scenarios where distinguishing between the absence of evidence and evidence of absence is critical.

Analysis of dataset dependency

Figure 2 visualizes the Bayes factors (BFs) for three LLMs—ChatGPT-3.5, DeepSeek-V3, and Llama3—on both the English and French versions of the CrowS-pairs dataset. The solid lines represent the BFs for the English CrowS-pairs, while the dashed lines correspond to the French CrowS-pairs. The DeepSeek-V3 results are depicted in green lines with triangles, Llama3 in red lines with squares, and ChatGPT-3.5 in blue lines with dots. The horizontal dotted-dashed line at 1 indicates the threshold for no evidence of bias, with additional dotted horizontal lines at 10 and 0.1 representing strong evidence for and against bias, respectively. Overall, all three LLMs exhibit significant consistency across nine bias categories in both the English and French CrowS-pairs datasets. DeepSeek-V3 (green lines with triangles) demonstrates extremely strong bias ($BF > 1e + 14$) across all nine categories in both languages. ChatGPT-3.5 (blue lines with dots) shows strong bias ($BF > 10$)




Fig. 2. The visualization of Bayes factors (BFs) for ChatGPT-3.5, DeepSeek-V3, and Llama3 on both the English and French versions of CrowS-pairs. The solid lines represent the BFs for the English CrowS-pairs, while the dashed lines correspond to the French CrowS-pairs. The dotted-dashed horizontal line at 1 indicates the threshold for no evidence of bias, with additional dotted horizontal lines at 10 and 0.1 representing strong evidence for and against bias, respectively. The results highlight the influence of cultural and linguistic context on bias detection outcomes.

in the Race and Religion categories across both languages. Llama3 (red lines with squares) exhibits strong bias ($BF > 10$) in the Age, Gender, Race, and Socioeconomic Status categories in both languages.

There are some significant difference in BFs between the English and French CrowS-Pairs datasets. For instance, ChatGPT-3.5 presents extremely strong evidence ($BF \approx 100$) of bias in Age in French, but it only exhibits anecdotal evidence in English. Both ChatGPT-3.5 and Llama3 exhibit moderate evidence of fairness ($BF < 1/3$) in Appearance in French, but they show no evidence of fairness in Appearance in English, where Llama3 even presents strong evidence of bias ($BF > 10$) in this category in English.

The observed differences in bias detection results between the English and French versions of the CrowS-Pairs dataset can be attributed to both linguistic and cultural factors. While the two datasets share a common structure, the translation process and cultural adaptation introduce nuances that influence how biases are expressed and detected. For instance, certain stereotypes like Age, may be more prevalent or differently framed in French-speaking cultures compared to English-speaking ones, leading to variations in the measured BFs.

Bias definitions are inherently context-dependent and shaped by cultural norms and societal values. In the English CrowS-Pairs dataset, biases are often framed within the context of American cultural norms, whereas the French version reflects the cultural context of French-speaking communities. For instance, religious biases may manifest differently in predominantly Catholic French-speaking cultures compared to more religiously diverse English-speaking contexts, which leads to the variations in computed BFs.

The current linguistic materials (English and French) were created within the US-English and France-French cultural contexts respectively, thus the observed cultural differences may be confined to these specific cultural frameworks rather than generalizable to all English- and French-speaking contexts. Notably, the Nationality bias likely correlates with specific national affiliations — while French is spoken in Canada, attitudinal differences may exist between France and Québec due to distinct sociopolitical environments.

Bias detection of LLMs on BBQ Dataset

Table 4 illustrates the bias detection results of ChatGPT-3.5-Turbo and DeepSeek-V3 over nine bias categories using the BBQ dataset. For DeepSeek-V3, both the EBT and BF methods consistently indicate strong bias across all nine categories. In contrast, ChatGPT-3.5-Turbo exhibits strong bias in seven categories, with exceptions in Race and Sexual Orientation.

For the Race category, the p -value from the EBT ($2.57e-1$) is not statistically significant at the 0.05 level, and the BF value of $5.9e-2$ provides strong evidence for the null hypothesis (no bias). In the Sexual Orientation category, the p -value from the EBT is insignificant at the 0.01 level but significant at the 0.05 level. The BF value of 1.23 suggests anecdotal evidence for bias in this category.

By comparing the results in Tables 3 and 4, we observe that DeepSeek-V3 demonstrates consistent bias across all nine categories in both the CrowS-Pairs and BBQ datasets. In contrast, ChatGPT-3.5-Turbo exhibits varying bias patterns across the two datasets. Specifically, on the English CrowS-Pairs dataset, ChatGPT-3.5-Turbo shows strong evidence of bias in the categories of Nationality, Race, and Religion, as indicated by high Bayes factors (e.g., $BF > 10$). However, on the BBQ dataset, it demonstrates strong evidence of bias in all categories except Race and Sexual Orientation, where the Bayes factors suggest moderate to strong evidence for no bias. These differences in bias detection results may be attributed to the distinct design and focus of the two datasets. The CrowS-Pairs dataset is designed to measure stereotypical biases through sentence pairs, while the BBQ dataset evaluates bias through question-answering tasks in ambiguous contexts. This variation in dataset structure and context likely influences the model’s responses, leading to the observed discrepancies in bias detection.

Effects of temperature and sample size

To assess the impact of temperature and sample size on bias detection in ChatGPT-3.5-Turbo, we conducted an investigation by varying these parameters. We subsampled the Sexual Orientation data from the BBQ dataset and computed bias metrics (SS, EBT, and BF) under different sample sizes (5–100% of the full dataset) and temperature settings (0.0–2.0). The results are presented in Fig. 3 and Table 5.

Figure 3 demonstrates the BFs (left panel) and p -values (right panel) of ChatGPT-3.5-Turbo on Sexual Orientation data of BBQ under varying sample sizes and temperature parameters. From the left panel, as the

Method	Bias								
	Age	Disability	Gender	Nationa.	Phy. App.	Race	Religion	Sex. Ori.	Socioeco.
	(920)	(389)	(1418)	(770)	(394)	(1720)	(300)	(216)	(1716)
ChatGPT-3.5-Turbo									
SS	71.41%	59.13%	62.20%	58.05%	60.41%	51.40%	70.00%	57.87%	55.07%
EBT	1.36e-39**	3.74e-04**	3.34e-20**	8.93e-06**	4.23e-05**	2.57e-01	3.31e-12**	2.45e-02*	2.92e-05**
BF	2.55e+36	4.21e+01	1.04e+17	1.01e+03	3.33e+02	5.9e-02	3.46e+09	1.23e+00	2.06e+02
DeepSeek-V3									
SS	92.39%	89.04%	75.57%	94.66%	83.91%	72.97%	86.72%	79.59%	78.96%
EBT	7.51e-101**	1.74e-34**	3.7e-09**	1.11e-44**	3.06e-23**	7.63e-03**	3.98e-14**	3.85e-05**	2.12e-37**
BF	5.1e+97	5.89e+31	5.97e+06	9.10e+41	4.01e+20	1.04e+01	5.95e+11	1.37e+03	2.73e+34

Table 4. Results of bias detection methods for nine biase types in two LLMs (ChatGPT-3.5, and DeepSeek-V3) on the BBQ dataset. The setup of this table is the same as Table 3.




Fig. 3. The visualization depicts Bayes factors (left panel) and p -values (right panel) for ChatGPT-3.5-Turbo under different sample sizes and temperature parameters. Data were subsampled from the Sexual Orientation subset of the BBQ dataset. The blue solid, red solid, blue dashed, yellow dash-dotted, and purple dotted lines correspond to settings with temperature parameters of 0.0, 0.5, 1.0, 1.5, and 2.0, respectively. The black horizontal segment represents the 0.05 significance level for p -values.

sample size increases, the BFs become divergent and have shown no strong evidence for either hypothesis. By contrast, the right panel presents a clear decline of p -values for all five temperature parameters as the increase of sample size. When the sample size is larger than 172, the p -values are significant under 0.05 level for temperature 0, 1, and 1.5. Therefore, this figure illustrates that BFs are more robust to sample size than p -values from EBT.

Benefit II of Bayes factors—robustness to sample size

Unlike the EBT, which is highly sensitive to sample size, the BF provides stable and interpretable results even with limited data. This makes it particularly useful in scenarios where data availability is constrained.

Table 5 presents the performance of three bias detection metrics—SS, EBT, and BF—for ChatGPT-3.5-Turbo across varying temperature parameters and sample sizes. For the special case where temperature = 0, the model becomes fully deterministic, producing identical outputs for given inputs. The results indicate that temperature = 0 yields relatively stable SS and BF values, particularly for larger sample sizes. EBT achieves significance ($p < 0.05$) for larger sample sizes, while BF demonstrates anecdotal evidence, suggesting greater robustness compared to EBT. Consequently, temperature = 0 serves as a baseline for assessing the model's inherent biases without the influence of randomness.

Higher temperatures (≥ 1.0) introduce increased variability in the model's outputs. Key observations include: (1.) SS fluctuates between 50.0% and 66.67% across sample sizes and temperatures; (2.) EBT p -values exhibit high sensitivity to sample size and temperature, achieving significance ($p < 0.05$) at temperatures 1.0 and 1.5 for larger sample sizes (80% and 100%). However, at temperature 2.0, EBT shows no significance for any sample size due to heightened output randomness; (3.) BF values frequently indicate moderate evidence (underlined values) for the null hypothesis (no bias). At temperature 2.0, BF consistently shows moderate evidence of no bias (fairness) across all sample sizes except 5%, likely due to output randomness at high temperatures.

In summary, both the temperature and sample size have significant impact on the bias detection of LLMs like ChatGPT. The high temperature (e.g., 2.0) increases the variability of output, leading to moderate evidence for fairness, while the low temperature allows no randomness, showing the inherent bias of the LLMs.

Gender bias detection with three datasets

Gender bias constitutes one of the most extensively studied forms of bias, attributable to its pervasive influence across diverse sectors such as employment, education, and social dynamics. The availability of numerous open-source datasets has facilitated the measurement of gender biases in LLMs. In this study, we demonstrate the application of BF to assess gender bias in two prominent models: ChatGPT-3.5-Turbo and DeepSeek-V3. Our analysis leverages three distinct datasets: Winogender, as well as the CrowS-Pairs datasets in both English (CrowS(EN)) and French (CrowS(FR)) languages.

Table 6 displays the results of three methods for gender bias in ChatGPT-3.5-Turbo and DeepSeek-V3. For ChatGPT-3.5-Turbo, the p -values from the EBT are statistically insignificant (e.g., $3.33\text{e-}1$ for Winogender, $4.02\text{e-}1$ for CrowS-Pairs English, and $1.18\text{e-}1$ for CrowS-Pairs French), indicating that we cannot reject the

Percentage	5%	10%	20%	50%	80%	100%
Temperature = 0.0						
SS	50.0%	50.0%	57.58%	57.95%	59.03%	58.79%
EBT	1.e+00	1.e+00	4.87e-01	1.65e-01	3.69e-02*	2.13e-02*
BF	3.69e-01	<u>2.70e-01</u>	<u>3.09e-01</u>	4.00e-01	1.08e+00	1.54e+00
Temperature = 0.5						
SS	66.67%	52.63%	51.35%	58.59%	55.19%	55.38%
EBT	5.08e-01	1.e+00	1.e+00	1.07e-01	2.27e-01	1.52e-01
BF	6.1e-01	<u>2.84e-01</u>	<u>2.05e-01</u>	5.34e-01	<u>2.3e-01</u>	<u>2.76e-01</u>
Temperature = 1.0						
SS	50.0%	57.14%	65.12%	53.70%	58.14%	57.87%
EBT	1.e+00	6.64e-01	6.6e-02	5.01e-01	3.92e-02*	2.45e-02*
BF	3.69e-01	3.24e-01	1.32e+00	<u>1.61e-01</u>	9.26e-01	1.23e+00
Temperature = 1.5						
SS	62.5%	50.0%	57.50%	60.58%	59.51%	57.07%
EBT	7.27e-01	1.e+00	4.3e-01	3.9e-02*	1.85e-02*	5.00e-02*
BF	5.08e-01	<u>2.70e-01</u>	<u>3.02e-01</u>	1.24e+00	1.86e+00	6.76e-01
Temperature = 2.0						
SS	66.67%	50.0%	57.89%	54.64%	51.66%	53.97%
EBT	5.08e-01	1.e+00	4.18e-01	4.17e-01	7.45e-01	3.09e-01
BF	6.1e-01	<u>2.84e-01</u>	<u>3.17e-01</u>	<u>1.91e-01</u>	<u>1.10e-01</u>	<u>1.64e-01</u>

Table 5. The robustness of bias detection metrics under varying sample sizes (ranging from 5 to 100% of the full dataset) and temperature parameters (five values from 0.0 to 2.0). The samples are drawn from the BBQ dataset under the bias category Sexual Orientation, and the model used here is **ChatGPT-3.5-Turbo**. As in Table 3, the *p*-values from the EBT followed by a single star (*) indicate it is significant at the 0.05 significance level. For BF, underlined values show moderate evidence for one of the two hypotheses ($1/10 < BF_{10} < 1/3$ or $3 < BF_{10} < 10$).

Method	Data					
	Winogender	CrowS(EN)	CrowS(FR)	Winogender	CrowS(EN)	CrowS(FR)
		ChatGPT-3.5			DeepSeek-V3	
SS	53.30%	52.50%	54.52%	85.41%	93.44%	96.67%
EBT	3.33e-1	4.02e-1	1.18e-1	1.97e-30**	4.09e-64**	1.96e-22**
BF	<u>1.37e-1</u>	<u>1.04e-1</u>	<u>2.58e-1</u>	5.03e+27	1.61e+61	1.16e+20

Table 6. Detecting gender bias in ChatGPT-3.5-Turbo and DeepSeek-v3 on three datasets: CrowS-Pairs in both English and French, and Winogender.

null hypothesis (H_0) of no bias. However, the *p*-value alone does not quantify the degree to which the data supports H_0 . In contrast, the BF provides a more nuanced interpretation. The BF values for ChatGPT-3.5-Turbo (e.g., 1.37e-1 for Winogender, 1.04e-1 for CrowS-Pairs English, and 2.58e-1 for CrowS-Pairs French) indicate **moderate to strong evidence** in favor of H_0 .

For **DeepSeek-V3**, the EBT results are highly significant (e.g., 1.97e-30 for Winogender, 4.09e-64 for CrowS-Pairs English, and 1.96e-22 for CrowS-Pairs French), providing strong evidence to reject H_0 . The corresponding BF values (e.g., 5.03e+27, 1.61e+61, and 1.16e+20) further confirm overwhelming evidence against H_0 , aligning with the EBT results.

Conclusion

In this paper, we formulate the bias detection of LMs as a hypothesis testing problem, and propose to utilize Bayes factors to quantify relative evidence for both competing hypotheses. Bayes factor has benefits over classical tests when the *p*-value greater than the predefined significance level, but it is rarely found in the natural language processing literature. Therefore, our work promotes the application of Bayes factors. We demonstrate the application of our framework to mainstream LLMs by leveraging three datasets as testbeds and illustrate the benefits of Bayes factors. This work represents a significant step toward developing fairer and more equitable language technologies, paving the way for responsible artificial intelligence deployment in real-world applications.

Data availability

The datasets employed in the preparation of this paper are publicly accessible via the following URLs: BBQ dataset: <https://github.com/nyu-ml/BQ>; CrowS-Pairs: <https://gitlab.inria.fr/french-crows-pairs/acl-2022-paper-data-and-code>; Winogender: <https://github.com/rudinger/winogender-schemas>.

Code availability

The code implemented in this study is available for public access at the following URL: https://github.com/shiji-ng001/bayes_factor_bias_detection.

Received: 2 February 2025; Accepted: 24 March 2025

Published online: 11 April 2025

References

- Gu, Y. *et al.* Eva2. 0: Investigating open-domain Chinese dialogue systems with large-scale pre-training. *arXiv preprint arXiv:2203.09313* (2022).
- Zhang, Y. *et al.* Dialogpt: Large-scale generative pre-training for conversational response generation. in *Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics: System Demonstrations*, 270–278 (2020).
- Radford, A., Narasimhan, K., Salimans, T. & Sutskever, I. Improving language understanding by generative pre-training (2018).
- Bao, S., He, H., Wang, F., Wu, H. & Wang, H. Plato: Pre-trained dialogue generation model with discrete latent variable. in *Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics*, 85–96 (2020).
- Nangia, N., Vania, C., Bhalerao, R. & Bowman, S. Crows-pairs: A challenge dataset for measuring social biases in masked language models. In *Proceedings of the 2020 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, 1953–1967 (2020).
- Bolukbasi, T., Chang, K.-W., Zou, J. Y., Saligrama, V. & Kalai, A. T. Man is to computer programmer as woman is to homemaker? Debiasing word embeddings. *Adv. Neural Inform. Process. Syst.* **29** (2016).
- Hutson, M. Robo-writers: The rise and risks of language-generating ai. *Nature* **591**(7848), 22–25 (2021).
- Brown, T. *et al.* Language models are few-shot learners. *Adv. Neural Inform. Process. Syst.* **33**, 1877–1901 (2020).
- Caliskan, A., Bryson, J. J. & Narayanan, A. Semantics derived automatically from language corpora contain human-like biases. *Science* **356**, 183–186 (2017).
- Davidson, T., Bhattacharya, D. & Weber, I. Racial bias in hate speech and abusive language detection datasets. In *Proceedings of the Third Workshop on Abusive Language Online*, 25–35, doi:10.18653/v1/W19-3504 (Association for Computational Linguistics, Florence, Italy, 2019).
- Delobelle, P., Tokpo, E., Calders, T. & Berendt, B. Measuring fairness with biased rulers: A comparative study on bias metrics for pre-trained language models. In *Proceedings of the 2022 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies*, 1693–1706. <https://doi.org/10.18653/v1/2022.naacl-main.122> (Association for Computational Linguistics, Seattle, United States, 2022).
- Kurita, K., Vyas, N., Pareek, A., Black, A. W. & Tsvetkov, Y. Measuring bias in contextualized word representations. In *Proceedings of the First Workshop on Gender Bias in Natural Language Processing*, 166–172. <https://doi.org/10.18653/v1/W19-3823> (Association for Computational Linguistics, Florence, Italy, 2019).
- Gallegos, I. O. *et al.* Bias and fairness in large language models: A survey. *Computational Linguistics* (2024).
- Parrish, A. *et al.* Bbq: A hand-built bias benchmark for question answering. In *Findings Assoc. Comput. Linguistics ACL 2022*, 2086–2105 (2022).
- Zhao, J., Wang, T., Yatskar, M., Ordonez, V. & Chang, K.-W. Gender bias in coreference resolution: Evaluation and debiasing methods. In *Proceedings of the 2018 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies, Volume 2 (Short Papers)*, 15–20 (2018).
- Dror, R., Baumer, G., Shlomov, S. & Reichart, R. The hitchhiker's guide to testing statistical significance in natural language processing. In *Proceedings of the 56th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, 1383–1392, <https://doi.org/10.18653/v1/P18-1128> (Association for Computational Linguistics, Melbourne, Australia, 2018).
- Zhiltsova, A., Caton, S. & Mulway, C. Mitigation of unintended biases against non-native english texts in sentiment analysis. In *Proceedings for the 27th AIAI Irish Conference on Artificial Intelligence and Cognitive Science, Galway, Ireland, December 5-6, 2019*, vol. 2563 of *CEUR Workshop Proceedings*, 317–328 (CEUR-WS.org, 2019).
- Dienes, Z. How Bayes factors change scientific practice. *J. Mat. Psychol.* **72**, 78–89 (2016).
- Heck, D. W. *et al.* A review of applications of the bayes factor in psychological research. *Psychol. Methods* **28**, 558 (2023).
- Qian, Y., Muaz, U., Zhang, B. & Hyun, J. W. Reducing gender bias in word-level language models with a gender-equalizing loss function. In *Proceedings of the 57th Annual Meeting of the Association for Computational Linguistics: Student Research Workshop*, 223–228 (2019).
- Yeo, C. & Chen, A. Defining and evaluating fair natural language generation. *arXiv preprint arXiv:2008.01548* (2020).
- Liu, R., Jia, C., Wei, J., Xu, G. & Vosoughi, S. Quantifying and alleviating political bias in language models. *Artif. Intell.* **304**, 103654. <https://doi.org/10.1016/j.artint.2021.103654> (2022).
- Guo, W. & Caliskan, A. Detecting emergent intersectional biases: Contextualized word embeddings contain a distribution of human-like biases. In *Proceedings of the AAAI/ACM Conference on AI, Ethics, and Society, AIES*, 122–133 (2021).
- Webster, K. *et al.* Measuring and reducing gendered correlations in pre-trained models. *CoRR arXiv:abs/2010.06032* (2020).
- Meade, N., Poole-Day, E. & Reddy, S. An empirical survey of the effectiveness of debiasing techniques for pre-trained language models. In Muresan, S., Nakov, P. & Villavicencio, A. (eds.) *Proceedings of the 60th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, 1878–1898, <https://doi.org/10.18653/v1/2022.acl-long.132> (Association for Computational Linguistics, Dublin, Ireland, 2022).
- Rudinger, R., Naradowsky, J., Leonard, B. & Durme, B. V. Gender bias in coreference resolution. In *Proceedings of the 2018 Conference of the North American Chapter of the Association for Computational Linguistics, NAACL*, 8–14 (2018).
- Dev, S., Li, T., Phillips, J. M. & Srikumar, V. On measuring and mitigating biased inferences of word embeddings. In *Proceedings of the 34th Association for the Advancement of Artificial Intelligence, AAAI*, 7659–7666 (2020).
- Nadeem, M., Bethke, A. & Reddy, S. Stereoset: Measuring stereotypical bias in pretrained language models. In *Proceedings of the 59th Annual Meeting of the Association for Computational Linguistics and the 11th International Joint Conference on Natural Language Processing (Volume 1: Long Papers)*, 5356–5371 (2021).
- Dinan, E. *et al.* Queens are powerful too: Mitigating gender bias in dialogue generation. In *Proceedings of the 2020 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, 8173–8188 (2020).
- Liu, H. *et al.* Mitigating gender bias for neural dialogue generation with adversarial learning. In *Proceedings of the 2020 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, 893–903 (2020).
- Barikeri, S., Lauscher, A., Vulić, I. & Glavaš, G. Redditbias: A real-world resource for bias evaluation and debiasing of conversational language models. In *Proceedings of the 59th Annual Meeting of the Association for Computational Linguistics and the 11th International Joint Conference on Natural Language Processing (Volume 1: Long Papers)*, 1941–1955 (2021).

32. Cheng, P., Hao, W., Yuan, S., Si, S. & Carin, L. Fairfil: Contrastive neural debiasing method for pretrained text encoders. In *International Conference on Learning Representations* (2021).
33. Cuddy, A. J. et al. Stereotype content model across cultures: Towards universal similarities and some differences. *Br. J. Social Psychol.* **48**, 1–33 (2009).
34. Fiske, S. T. Prejudices in cultural contexts: Shared stereotypes (gender, age) versus variable stereotypes (race, ethnicity, religion). *Perspect. Psychol. Sci.* **12**, 791–799 (2017).
35. Hinton, P. Implicit stereotypes and the predictive brain: Cognition and culture in “biased” person perception. *Palgrave Commun.* **3**, 1–9 (2017).
36. Névéol, A., Dupont, Y., Bezançon, J. & Fort, K. French CrowS-pairs: Extending a challenge dataset for measuring social bias in masked language models to a language other than English. In *Proceedings of the 60th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, 8521–8531, <https://doi.org/10.18653/v1/2022.acl-long.583> (Association for Computational Linguistics, Dublin, Ireland, 2022).
37. Radford, A. et al. Language models are unsupervised multitask learners. *OpenAI Blog*. **1**, 9 (2019).
38. Bhatia, N. B. Changes in gender stereotypes over time: A computational analysis. *Psychol. Women Quart.* **45**, 106–125. <https://doi.org/10.1177/0361684320977178> (2021).
39. Andraszewicz, S. et al. An introduction to Bayesian hypothesis testing for management research. *J. Manag.* **41**, 521–543 (2015).
40. Geisser, S. On prior distributions for binary trials. *Am. Statistician* **38**, 244–247. <https://doi.org/10.1080/00031305.1984.10483216> (1984).

Acknowledgements

This work was supported by the Humanities and Social Sciences Research Youth Foundation of Ministry of Education of China under Grant Number 24YJCZH252; Special Project on Artificial Intelligence-Driven Research Paradigm Reform and Discipline Leapfrog Development Empowered by Shanghai Municipal Education Commission (Project Name: Research on the Development and Applications of Performance Evaluation Datasets for Large Language Models). We would like to thank Professor Kaibao Hu, and other members in the Artificial Intelligence and Humanities & Social Sciences Interdisciplinary Research Team of Shanghai International Studies University for many insightful discussions about the detection of social biases in this work, as well as the reviewers for their helpful comments.

Author contributions

S.S. conceived and conducted the experiments, X.J. and S.S. analysed the results and wrote the manuscript. All authors reviewed the manuscript.

Declarations

Competing interests

The authors declare no competing interests.

Approval for human experiments

Given that this study did not involve human or animal subjects, and the APIs for ChatGPT, DeepSeek-V3, and Llama-3.1 are publicly available and accessible online, formal approval from an ethical review committee was not required. All data utilized in this research were obtained from open-access sources, ensuring compliance with ethical standards for non-interventional studies.

Additional information

Correspondence and requests for materials should be addressed to X.J.

Reprints and permissions information is available at www.nature.com/reprints.

Publisher’s note Springer Nature remains neutral with regard to jurisdictional claims in published maps and institutional affiliations.

Open Access This article is licensed under a Creative Commons Attribution-NonCommercial-NoDerivatives 4.0 International License, which permits any non-commercial use, sharing, distribution and reproduction in any medium or format, as long as you give appropriate credit to the original author(s) and the source, provide a link to the Creative Commons licence, and indicate if you modified the licensed material. You do not have permission under this licence to share adapted material derived from this article or parts of it. The images or other third party material in this article are included in the article’s Creative Commons licence, unless indicated otherwise in a credit line to the material. If material is not included in the article’s Creative Commons licence and your intended use is not permitted by statutory regulation or exceeds the permitted use, you will need to obtain permission directly from the copyright holder. To view a copy of this licence, visit <http://creativecommons.org/licenses/by-nc-nd/4.0/>.

© The Author(s) 2025